

# parameter learning $\theta$ gradient descent basics

## gradient descent

the use of gradient descent as a mechanism for minimizing the function of  $J(\theta_1, \theta_0)$

note: for brevity, the functions  $J(\theta_1, \theta_0)$  are shortened throughout the document

$$J(\theta_0, \theta_1, \theta_2, \dots, \theta_n) \rightarrow J(\theta_1, \theta_0)$$


$$\min_{\theta_0, \theta_1, \theta_2, \dots, \theta_n} J(\theta_0, \theta_1, \theta_2, \dots, \theta_n) \rightarrow \min_{\theta_0, \theta_1} J(\theta_1, \theta_0)$$

the process begins by initializing some value of parameters  $\theta_0, \theta_1$  (e.g.  $\theta_0 = 0, \theta_1 = 0$ )

the values  $\theta_0, \theta_1$  will then change through reiterations to reduce the value of  $J(\theta_0, \theta_1)$

the reiterations will run continuously with the intended arrival at a **minimum**

illustrated below, the initialized values will be plotted against the function:



in narrative explanation, gradient descent begins at the points initialized and scans the contour in a 360°-like fashion and chose a direction the descends the counter in the most rapid route. the initialized values will update to the new point on the contour and repeat the process until the functions determines it has reached a **local minimum**

a unique property of gradient descent occurs dependent upon the initialized values of parameters . different initialization values can result in a disjoint **local minimum**

the formal gradient descent algorithm is as follows:

repeat until convergence {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) \quad (\text{for } j = 0 \text{ and } j = 1)$$

}

the  $\alpha$  term determines the magnitude of each step in gradient descent

the term  $\frac{\partial}{\partial \theta_j}$  is the partial derivative active in gradient descent

the gradient descent algorithm functions as a simultaneous update of  $\theta_0, \theta_1$ :

correct simultaneous update:

$$\text{temp0} := \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$$

$$\text{temp1} := \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$$

$$\theta_0 := \text{temp0}$$

$$\theta_1 := \text{temp1}$$

incorrect simultaneous update:

$$\text{temp0} := \theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$$

$$\theta_0 := \text{temp0}$$

$$\text{temp1} := \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$$

$$\theta_1 := \text{temp1}$$

the incorrect example of simultaneous update displays the use of the updated value of  $\theta_0$  in the computation of updating  $\theta_1$

Suppose  $\theta_0 = 1, \theta_1 = 2$ , and we simultaneously update  $\theta_0$  and  $\theta_1$  using the rule:

$\theta_j := \theta_j + \sqrt{\theta_0 \theta_1}$  (for  $j = 0$  and  $j = 1$ ) What are the resulting values of  $\theta_0$  and  $\theta_1$ ?


- ☐  $\theta_0 = 1, \theta_1 = 2$
- ☒  $\theta_0 = 1 + \sqrt{2}, \theta_1 = 2 + \sqrt{2}$

Correct Response

- ☐  $\theta_0 = 2 + \sqrt{2}, \theta_1 = 1 + \sqrt{2}$
- ☐  $\theta_0 = 1 + \sqrt{2}, \theta_1 = 2 + \sqrt{(1 + \sqrt{2}) \cdot 2}$

## gradient descent intuition


the function of the **derivative** in the gradient descent function takes the tangent slope of the line plotted against  $\theta_1$  and multiplies it against learning rate  $\alpha$  and cost function  $J(\theta_1)$ ; the functions then applies the direction and magnitude into a step



$$(\theta_1 \in \mathbb{R})$$

$$\theta_1 := \theta_1 - \alpha \frac{d}{d\theta_1} J(\theta_1)$$
$$\geq 0$$

$$\theta_1 := \theta_1 - \alpha(\text{positive number})$$



$$\theta_1 := \theta_1 - \alpha \frac{d}{d\theta_1} J(\theta_1)$$
$$\leq 0$$


$$\theta_1 := \theta_1 - \alpha(\text{negative number})$$

the initialized values in the gradient descent function will determine whether the tangent slope of the derivative will require a **positive** or **negative** number to make a step towards finding a local minimum

with reference to the behavior of the assigned learning rate  $\alpha$ :



$$\theta_1 := \theta_1 - \alpha \frac{d}{d\theta_1} J(\theta_1)$$

if the learning rate  $\alpha$  is too **small**, gradient descent can process slowly and take substantial time and resources converging to a local optima (minimum)




$$\theta_1 := \theta_1 - \alpha \frac{d}{d\theta_1} J(\theta_1)$$

if the learning rate  $\alpha$  is too **large**, gradient descent can overshoot the local optima (minimum); it might equally fail to converge, or diverge entirely from obtaining a local optima (minimum)



if  $\theta_1$  is already at a local optimum of  $J(\theta_1)$ , a single additional step gradient descent  $\theta_1 := \theta_1 - \alpha \frac{d}{d\theta_1} J(\theta_1)$  will leave  $\theta_1$  unchanged at local minimum, the tangent slope of the derivative = 0 and thus will not affect  $\theta_1$



if the learning rate alpha  $\alpha$  is fixed, gradient descent can still converge to a local minimum

$$\theta_1 := \theta_1 - \alpha \frac{d}{d\theta_1} J(\theta_1)$$

after each simultaneous update of gradient descent, the algorithm autonomously takes smaller steps due to a decreasing derivative computation with each iteration

## gradient descent for linear regression

applying gradient descent to minimize the linear regression cost function:

gradient descent algorithm

repeat until convergence {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$$

(for  $j = 0$  and  $j = 1$ )

}

linear regression model

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$$\min_{\theta_0, \theta_1} J(\theta_0, \theta_1)$$

determining the function of the derivative term  $\frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$ :

$$\begin{aligned} \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) &= \frac{\partial}{\partial \theta_j} \times \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 \\ &= \frac{\partial}{\partial \theta_j} \times \frac{1}{2m} \sum_{i=1}^m (\theta_0 + \theta_1 x^{(i)} - y^{(i)})^2 \end{aligned}$$

determining the partial derivatives in both cases of  $\theta_0, \theta_1$

the partial derivatives are simply the slopes of the cost function  $J(\theta_0, \theta_1)$

$$\theta_0 \rightarrow j = 0: \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$$\theta_1 \rightarrow j = 1: \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 x^{(i)}$$

gradient descent algorithm for linear regression:

repeat until convergence {

$$\theta_0 := \theta_0 - a \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})$$

$$\theta_1 := \theta_1 - a \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x^{(i)}$$

}

$$\frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$$


update  $\theta_0$  and  $\theta_1$  simultaneously

$$\frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$$

the cost function for linear regression will always be represented as a **convex function**


"**batch**" gradient descent is a common term in machine learning to refer to each step of gradient descent using all of the training examples in the dataset

convex functions do not consist of any local optima, there is now only an available **global optima** (global minimum) of the convex function



$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )



$$J(\theta_0, \theta_1)$$

(function of the parameters  $\theta_0, \theta_1$ )

