

logistic regression ↗ model

cost function

given a training set: $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$

with m examples denoted as $x \in \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_n \end{bmatrix}$; where $x_0 = 1, y \in \{0,1\}$

and a hypothesis: $h_{\theta}(x) = \frac{1}{1+e^{-\theta^T x}}$

the parameters θ can be determined as follows:


linear regression:

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2 \rightarrow \text{cost}(h_{\theta}(x^{(i)}), y^{(i)})$$

$$\text{cost}(h_{\theta}(x^{(i)}), y^{(i)}) = \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$


the simplified cost functions serves purpose as to measure the cost the training set will have to pay when $h_{\theta}(x^{(i)})$ is predicted in terms of actual $y^{(i)}$: being $\frac{1}{2}$ the squared error

the Linear Regression Cost Function cannot be applied to logistic regression as it would not create a Convex Function; Gradient Descent would fail to find the global minimum:



instead, logistic regression: $\text{cost}(h_{\theta}(x^{(i)}), y^{(i)}) = \begin{cases} -\log(h_{\theta}(x^{(i)})) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x^{(i)})) & \text{if } y = 0 \end{cases}$

the cost function is illustrated in each class as follows:




cost = 0 if $y = 1, h_{\theta}(x) = 1$

but as $h_{\theta}(x) \rightarrow 0$

cost $\rightarrow \infty$

the illustration (left) captures the intuition if $h_{\theta}(x^{(i)}) = 0$, (predict $P(y = 1|x:\theta) = 0$), but $y = 1$, the learning algorithm will be heavily penalized with a significant measured cost



cost = 0 if $y = 0, h_{\theta}(x) = 0$

but as $h_{\theta}(x) \rightarrow 1$

cost $\rightarrow \infty$

the illustration (left) captures the intuition if $h_{\theta}(x^{(i)}) = 1$, (predict $P(y = 0|x:\theta) = 0$), but $y = 0$, the learning algorithm will be heavily penalized with a significant measured cost

In logistic regression, the cost function for our hypothesis outputting (predicting) $h_{\theta}(x)$ on a training example that has label $y \in \{0, 1\}$ is:

$$\text{cost}(h_{\theta}(x), y) = \begin{cases} -\log h_{\theta}(x) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$

Which of the following are true? Check all that apply.

☒ If $h_{\theta}(x) = y$, then $\text{cost}(h_{\theta}(x), y) = 0$ (for $y = 0$ and $y = 1$).

Correct Response

☒ If $y = 0$, then $\text{cost}(h_{\theta}(x), y) \rightarrow \infty$ as $h_{\theta}(x) \rightarrow 1$.

Correct Response

☐ If $y = 0$, then $\text{cost}(h_{\theta}(x), y) \rightarrow \infty$ as $h_{\theta}(x) \rightarrow 0$.

Correct Response

☒ Regardless of whether $y = 0$ or $y = 1$, if $h_{\theta}(x) = 0.5$, then $\text{cost}(h_{\theta}(x), y) > 0$.

Correct Response

simplified cost function and gradient descent

logistic regression:

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{cost}(h_{\theta}(x^{(i)}), y^{(i)})$$

$$\text{cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$

note: $y = 0$ or $y = 1$ always

explanation of simplifying the cost function $\text{cost}(h_{\theta}(x^{(i)}), y^{(i)})$ above:

$$\text{cost}(h_{\theta}(x^{(i)}), y^{(i)}) = -y \log(h_{\theta}(x)) - (1 - y) \log(1 - h_{\theta}(x))$$

if $y=1$: $\text{cost}(h_{\theta}(x), y) = -\log(h_{\theta}(x))$

if $y=0$: $\text{cost}(h_{\theta}(x), y) = -\log(1 - h_{\theta}(x))$

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{cost}(h_{\theta}(x^{(i)}), y^{(i)})$$

$$= \frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \right]$$

principle of likelihood maximization derived function (statistical rational behind cost function used)

to fit parameters θ :

the process of minimizing $J(\theta)$; $\min_{\theta} J(\theta)$ will provide the parameters θ

to make a prediction given a new x training example:

output $h_{\theta}(x) = \frac{1}{1+e^{-\theta^T x}}$; $P(y = 1|x; \theta)$

in order to minimize $J(\theta)$; $\min_{\theta} J(\theta)$ (gradient descent)

repeat until convergence {

$$\theta_j := \theta_j - a \frac{\partial}{\partial \theta_j} J(\theta) = \theta_j - \alpha \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

→ (simultaneously update all θ_j)

$$\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \vdots \\ \theta_n \end{bmatrix}$$

updating all values of parameter θ is best performed as a vectorized

note the identical cosmetic similarity between the logistic and linear regression update rule is attributed to the altered definition of $h_{\theta}(x)$ from $h_{\theta}(x) = \theta^T X$ to $h_{\theta}(x) = \frac{1}{1+e^{-\theta^T x}}$

Suppose you are running gradient descent to fit a logistic regression model with parameter $\theta \in \mathbb{R}^{n+1}$. Which of the following is a reasonable way to make sure the learning rate α is set properly and that gradient descent is running correctly?

- ☐ Plot $J(\theta) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$ as a function of the number of iterations (i.e. the horizontal axis is the iteration number) and make sure $J(\theta)$ is decreasing on every iteration.
- ☒ Plot $J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))]$ as a function of the number of iterations and make sure $J(\theta)$ is decreasing on every iteration.

Correct Response

- ☐ Plot $J(\theta)$ as a function of θ and make sure it is decreasing on every iteration.
- ☐ Plot $J(\theta)$ as a function of θ and make sure it is convex.

One iteration of gradient descent simultaneously performs these updates:

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_0^{(i)}$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_1^{(i)}$$

\vdots

$$\theta_n := \theta_n - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_n^{(i)}$$

We would like a vectorized implementation of the form $\theta := \theta - \alpha \delta$ (for some vector $\delta \in \mathbb{R}^{n+1}$).

What should the vectorized implementation be?

- ☒ $\theta := \theta - \alpha \frac{1}{m} \sum_{i=1}^m [(h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x^{(i)}]$

Correct Response

- ☐ $\theta := \theta - \alpha \frac{1}{m} [\sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})] \cdot x^{(i)}$
- ☐ $\theta := \theta - \alpha \frac{1}{m} x^{(i)} [\sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})]$

advanced optimization

optimization algorithm

cost function $J(\theta)$ wanting to $\min_{\theta} J(\theta)$

given θ , code is written to compute

$$J(\theta) \text{ and } \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

(for $j = 0, 1, \dots, n$)

gradient descent:

repeat {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

}

optimization algorithms options:

- .. gradient descent
- .. conjugate gradient
- .. bfgs
- .. l-bfgs

advantages:

- .. unnecessary to select α manually
- .. often faster than gradient descent

disadvantages:

- .. generally more complex

optimization through example:

goal to $\min_{\theta} J(\theta)$ assuming $\theta_1 = 5$, $\theta_2 = 5$:

$$\theta = \begin{bmatrix} \theta_1 \\ \theta_2 \end{bmatrix}$$

$$J(\theta) = (\theta_1 - 5)^2 + (\theta_2 - 5)^2$$

$$\frac{\partial}{\partial \theta_j} J(\theta) = 2(\theta_1 - 5)$$

$$\frac{\partial}{\partial \theta_j} J(\theta) = 2(\theta_2 - 5)$$

optimization in octave:

```
function [jVal, gradient]
    = costFunction(theta)
    jVal = (theta(1)-5)^2 + ...
          (theta(2)-5)^2
    gradient = zeros(2,1)
    gradient(1) = 2*(theta(1)-5);
    gradient(2) = 2*(theta(2)-5);
```

after implementing the cost function, call the advanced optimization function in octave:

```
options = optimset('GradObj', 'on', 'MaxIter', '100');
initialTheta = zeros(2,1);
[optTheta, functionVal, exitFlag]...
    = fminunc(@costFunction, initialTheta, options);
```

'GradObj' in reference to gradient objective set to 'on' in reference to the fact that a gradient is going to be provided to the algorithm. `initialTheta` initializes the parameters θ are at $\mathbf{0}$. the advanced function `fminunc` is called to compute `optTheta`, the learning rate α autonomously:

```
octave-3.2.4.exe:1> PS1('>> ')
>> cd 'C:\Users\ang\Desktop'
>>
>> options = optimset('GradObj','on', 'MaxIter', '100');
>> initialTheta = zeros(2,1)
initialTheta =

    0
    0

<ag> = fminunc(@costFunction, initialTheta, options)

<ag> = fminunc(@costFunction, initialTheta, options)
optTheta =

    5.0000
    5.0000

functionVal = 1.5777e-030
exitFlag = 1
>>
```

application of algorithm optimization to logistic regression

$\text{theta} = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_n \end{bmatrix}$ $\begin{matrix} \leftarrow \text{theta}(1) \\ \leftarrow \text{theta}(2) \\ \vdots \\ \leftarrow \text{theta}(n) \end{matrix}$ *note the definition of theta being indexed at 1 in octave as opposed to 0 in the*

```
function [jVal, gradient] = costFunction(theta)

jVal = [code to compute J(theta)];

gradient(1) = [code to compute  $\frac{\partial}{\partial \theta_0} J(\theta)$ ];

gradient(2) = [code to compute  $\frac{\partial}{\partial \theta_1} J(\theta)$ ];

:

gradient(n+1) = [code to compute  $\frac{\partial}{\partial \theta_n} J(\theta)$ ];
```

the above code requires a user to input code to compute both a cost function and relative gradients

Suppose you want to use an advanced optimization algorithm to minimize the cost function for logistic regression with parameters θ_0 and θ_1 . You write the following code:

```
function [jVal, gradient] = costFunction(theta)
    jVal = % code to compute J(theta)
    gradient(1) = CODE#1 % derivative for theta_0
    gradient(2) = CODE#2 % derivative for theta_1
```

What should CODE#1 and CODE#2 above compute?

- ☐ CODE#1 and CODE#2 should compute $J(\theta)$.
- ☐ CODE#1 should be theta(1) and CODE#2 should be theta(2).
- ☒ CODE#1 should compute $\frac{1}{m} \sum_{i=1}^m [(h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_0^{(i)}] (= \frac{\partial}{\partial \theta_0} J(\theta))$ and


CODE#2 should compute $\frac{1}{m} \sum_{i=1}^m [(h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x_1^{(i)}] (= \frac{\partial}{\partial \theta_1} J(\theta))$

Correct Response


multiclass classification

multiclass classification: one-vs-all

when classification problems have more than a binary classification of 0 or 1



this is possible by separating examples into individual binary classification problems



formally, multiclass classification trains a logistic regression classifier

$h_{\theta}^{(i)}(x)$ for each class i that maximizes $\max_i h_{\theta}^{(i)}(x)$

$$h_{\theta}^{(i)}(x) = P(y = i|x; \theta)$$

where $(i = 1, 2, 3)$