

# error metrics ~~with~~ handling skewed data

## error metrics for skewed classes

cancer classification example

train a logistic regression model  $h_{\theta}(x)$   $y = \begin{cases} 1 & \text{if cancer} \\ 0 & \text{otherwise} \end{cases}$

the test set results in 1% error (99% correct diagnoses')

only 50% of the patients contract cancer (the classes are skewed)

```
function y = predictCancer(x)    → 0.5% error
    y = 0; %ignore x!
return
```

→ 99.2% accuracy (0.8% error)      → 99.5% accuracy (0.5% error)

an alternative evaluation metric for the above problem

precision/recall

$y = 1$  in presence of rare class that is the target to detect

		actual class	
		1	0
predicted class	1	true positives (+)	false positives (+)
	0	false negatives (-)	true negatives (-)

precision

(of all patients where prediction = 1, what fraction actually has cancer?)

$$\frac{\text{true positives}}{\# \text{ predicted positives}} = \frac{\text{true positives}}{\text{true positives} + \text{false positives}}$$

recall

(of all patients that actually have cancer, what fraction was correctly detected as having cancer?)

$$\frac{\text{true positives}}{\# \text{ actual positives}} = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}}$$

## trading off precision and recall

logistic regression:  $0 \leq h_{\theta}(x) \leq 1$

predict 1 if  $h_{\theta}(x) \geq 0.5$  (default)

predict 0 if  $h_{\theta}(x) < 0.5$  (default)

suppose prediction of  $y = 1$  (cancer) only if very confident

predict 1 if  $h_{\theta}(x) \geq 0.7$  or  $0.9$

predict 0 if  $h_{\theta}(x) < 0.7$  or  $0.9$

→ the model will produce **higher precision** but **lower recall**


suppose avoidance to missing too many cases of cancer (avoid false negatives)

predict 1 if  $h_{\theta}(x) \geq 0.7$  or  $0.9$

predict 0 if  $h_{\theta}(x) < 0.7$  or  $0.9$

→ the model will produce **lower precision** but **higher recall**

in general terms: predict 1 if  $h_{\theta}(x) \geq$  **given threshold**



## F<sub>1</sub> score

### comparing precision/recall metrics

precision and recall illustrate relationships among the two. using these metrics to select algorithms requires performance to be measured as a whole. the table below displays:

	precision (p)	recall (r)	average	F <sub>1</sub> score
algorithm 1	0.5	0.4	<del>0.45</del>	0.444
algorithm 2	0.7	0.1	<del>0.4</del>	0.175
algorithm 3	0.02	1.0	<del>0.51</del>	0.0392

$$\text{average} = \frac{p+r}{2}$$

$$\text{F}_1 \text{ score} = 2 \frac{pr}{p+r}$$

the average of precision and recall dictate that the 3<sup>rd</sup> algorithm is most attractive. this is misleading because that model predicted  $y = 1$  for the entire sample (recall = 100%) and greatly skewed the average performance between precision and recall

the F<sub>1</sub> score applies a weight to the lower value, making it penalize the metric accordingly. the illustration above shows the first algorithm as the most effective and conversely the 3<sup>rd</sup> algorithm substantially worse than compared to the average metric.

if  $P = 0$  or  $R = 0 \rightarrow \text{F}_1 \text{ score} = 0$

if  $P = 1$  or  $R = 1 \rightarrow \text{F}_1 \text{ score} = 1$

### example

a logistic c regression classifier has been trained as follows:

predict  $y = 1$  if  $h_{\theta}(x) \geq \text{threshold}$

predict  $y = 0$  if  $h_{\theta}(x) < \text{threshold}$

different values for the threshold parameter will yield different values of precision (p) and recall (r). a reasonable way to determine the threshold would be to:

measure precision (p) and recall (r) on the **cross validation set** and choose the value of threshold which maximizes the  $\text{F}_1 \text{ score} = 2 \frac{pr}{p+r}$

## using large data sets


### data for machine learning

determining how much data to train on  
designing a high accuracy learning system

example classify between confusable words

{to, two, too} {then, than} {there, their}

→ "For breakfast I ate two eggs."



the algorithms used to predict the word that belongs in the sentence above:

- perceptron (logistic regression)
- winnow
- memory-based
- naïve bayes

"it is not who has the best algorithm that wins, but who has the most data"

### large data rationale

assuming feature  $x \in \mathbb{R}^{n+1}$  has sufficient information to predict  $y$  accurately

example: "For breakfast I ate two eggs."

counterexample" "Predict housing price from only the size (feet<sup>2</sup>) as a feature."

useful test: given the input  $x$ , can a human expert confidently predict  $y$ ?

using a learning algorithm with many parameters (e.g. logistic regression/linear regression with many features; neural network with many hidden units).

→ low bias algorithms

$J_{train}(\theta)$  will be small

using a very large training set (unlikely to overfit)

→ low variance algorithms

$J_{train}(\theta) \approx J_{test}(\theta)$

$J_{test}(\theta)$  will be small

large training sets can help improve a learning algorithm's performance, except:  
when the features  $x$  do not contain enough information to predict  $y$ , regardless if the chosen algorithm is simple (logistic regression) or complex (neural networks)

# Machine Learning System Design

1. You are working on a spam classification system using regularized logistic regression. "Spam" is a positive class ( $y = 1$ ) and "not spam" is the negative class ( $y = 0$ ). You have trained your classifier and there are  $m = 1000$  examples in the cross-validation set. The chart of predicted class vs. actual class is:

	Actual Class: 1	Actual Class: 0
Predicted Class: 1	85	890
Predicted Class: 0	15	10

For reference:

- Accuracy = (true positives + true negatives) / (total examples)
- Precision = (true positives) / (true positives + false positives)
- Recall = (true positives) / (true positives + false negatives)
- $F1$  score =  $(2 * \text{precision} * \text{recall}) / (\text{precision} + \text{recall})$

What is the classifier's  $F1$  score (as a value from 0 to 1)? 0.16

2. Suppose a massive dataset is available for training a learning algorithm. Training on a lot of data is likely to give good performance when:

The features  $x$  contain sufficient information to predict  $y$  accurately. (For example, one way to verify this is if a human expert on the domain can confidently predict  $y$  when given only  $x$ ).

We train a learning algorithm with a large number of parameters (that is able to learn/represent fairly complex functions).

3. Suppose you have trained a logistic regression classifier which is outputting  $h\theta(x)$ .

Currently, you predict 1 if  $h\theta(x) \geq \text{threshold}$ , and predict 0 if  $h\theta(x) < \text{threshold}$ , where currently the threshold is set to 0.5.

Suppose you **increase** the threshold to 0.9:

The classifier is likely to now have higher precision.

4. Suppose you are working on a spam classifier, where spam emails are positive examples ( $y=1$ ) and non-spam emails are negative examples ( $y=0$ ). You have a training set of emails in which 99% of the emails are non-spam and the other 1% is spam:

If you always predict non-spam (output  $y=0$ ), the classifier will have accuracy of 99%.

If you always predict non-spam (output  $y=0$ ), your classifier will have 99% accuracy on the training set, and it will likely perform similarly on the cross validation set.

5. Using a **very large** training set makes it unlikely a model to overfit the training data.

On skewed datasets (e.g., when there are more positive examples than negative examples), accuracy is not a good measure of performance and should instead use  $F1$  score based on the precision and recall.